

Front Doors — presenter guide

What to ask it with. Runs after The Stranger, which covered how to phrase the request. About 12 minutes, or 4 if you are behind. Product claims checked 8 September 2026.

The one sentence

The surface you use decides what the tool can reach — and the further it reaches, the more of its work you have to check.

Everything on the page serves that sentence. If you only get one thing across, get that, and specifically the second half, because the first half is intuitive and the second is not.

Why this follows The Stranger

The Stranger taught the room that how you ask changes what you get. The obvious next conclusion — and the wrong one — is that prompting is the whole skill. This module is the correction: a perfectly phrased request to a browser tab still cannot rename a file on your laptop. Phrasing is one variable; the door is the other.

Running it — the 12-minute version

Act I • the four doors 4 MIN

Put the grid up and say nothing for a beat. Let them see four columns and read the four questions down the left. Then:

- **Open one cell, not six.** The best single cell in the module is **door 1, row 2** — a chat window answers in the conversation. Then jump straight to **door 1 on the “rename four hundred files” job** in Act II if you want the laugh, or stay for row 4.
- **Then open door 2, row 4** — accept or reject, you get a diff. This is the cell that surprises people, because it is the least glamorous door and it has the best review story in the table.
- **Name the bars honestly.** If somebody says the staircase looks suspiciously neat, agree with them: rows 1 to 3 rise because the doors were chosen as a ladder. Point at row 4 and say that is the row that was not chosen, it came out that way.

“Same model behind all four of these. What changes is how far it can reach — and every step of reach you hand it, you get back as review.”

Act II • pick a job 6 MIN

This is the act that lands. Do not walk all 7 jobs. Ask the room to pick, or pick two yourself:

- **“Rename four hundred scanned files”** — the clearest floor in the module. Two doors say *wrong door* outright. Nobody argues with it, and it makes the point that better prompting cannot buy reach.
- **“Reformat a syllabus into the department’s template”** — the one where the answer is the boring door, the agent already inside Docs. Most useful answer of the day for most of the room.
- **Check door 2 actually works on your account before you present it.** Murray State runs Google Workspace, which is why this module leads door 2 with Docs and Sheets rather than Word and Excel — but running Workspace is *not* the same as having Gemini switched on in it. That is licensed and enabled by an administrator, per account. Open a Doc and look for the pane the morning of the talk. If it is not there, say so from the stage and point at door 1: “this is the door I would want you to have, and you may need to ask IT for it” is a perfectly good beat, and far better than a demo that will not open.
- **Read the last line of every card out loud.** “You hand over”. That line is the whole argument. A capability chart never has it.

The green box at the top is computed, not written: it is the smallest-numbered door whose verdict is *does the whole job*. Say “smallest” deliberately — the module is not asking what the most capable tool is.

“Notice what I did not ask. Not which is best. Which is smallest and still finishes the job.”

Act III · the receipts 2 MIN

Do not read this act. Show two things and move on:

- **“Agent Mode”** → **“Edit with Copilot”**, with Microsoft’s own sentence under it. Same capability, new name, inside six months.
- **“ChatGPT agent”** → **retired**. Point out that OpenAI’s help page still explains how to use the thing it opens by telling you is gone. That is the state of the landscape, from the vendor.

If the room is engaged, the third entry is worth ten seconds: the brief this module was built from called the product “GPT Work”, which has never existed. Somebody who does this for a living got it wrong six weeks out of date. That is the argument for the date at the top of the page, and it is more persuasive as a confession than as advice.

The 4-minute cut

1. Grid up. Read the four questions down the left. **40 seconds.**
2. Open **door 2, row 4** — accept or reject, you get a diff. **50 seconds.**
3. Act II, the file-renaming job. Two doors say *wrong door*. **90 seconds.**
4. Act III, the two renames. **40 seconds.** Then the one sentence, and stop.

Questions you will get

“So which one should I buy?”

The module deliberately does not answer this and neither should you. The honest reply is: almost certainly none, because your institution already licenses one or two of these and that decides it. Send them to IT. It is not a dodge — the institutional answer genuinely dominates the vendor comparison here.

“Is the expensive one better?”

No, and row 4 is why. More reach is more review. For most of what a faculty member does in a week, doors 1 and 2 are the right answer — and of the 7 jobs in Act II, 3 are finished completely by a chat window or by the agent inside the word processor.

“Can it see my students’ work?”

Only if you give it access, and whether you may give it access is a question for your registrar and your IT department rather than for a vendor page. Two of the jobs in Act II carry that caution on screen for exactly this moment; open one and read it rather than improvising a policy answer.

“That product changed — your page is wrong”

“Good — that is the module working. Which door was it? Because the door is still there.”

Then open Act III and show that the page predicted exactly this, and lists what it expects to rot first. Do not defend the name. This is the strongest moment available to you and it only works if you take the correction gratefully.

“How do you know any of this?”

Open **Settings** → **Run the self-test** in front of them. It checks that every claim resolves to a declared source, that no source is cited and never used, that no price and no vendor ranking leaked into anything the page can display, and that the review row genuinely does not rise the way the other three do. Then tap any cell and follow the link.

What the numbers on the page mean

FIGURE	WHAT IT IS
16 cells	4 doors × 4 axes. Every one opens a reason, a label and a source.
13 verified	Cells where a page belonging to the company that makes the product was opened on the check date and says it.
3 reasoned	The ordering judgements. Labelled, because no vendor states them.
28 judgements	7 jobs × 4 doors in Act II. All reasoned from Act I, and each names the axis that decided it.
15 sources	All first-party. Two of them are OpenAI pages that refuse automated retrieval and were read in an ordinary browser.
9 things cut	Researched during the build and deliberately left off, each with the reason.

The shipped test suite fails the build if any of those figures disagrees with the code, this guide, README.md or the manifest. If you edit a number here, run **node src/playtest.test.js** before you present.

What is deliberately not on the page

- **No prices.** Not one figure. They moved more than once in the last year on all four doors and a wrong number in front of faculty costs more than a missing one. Figures that sources did state are in README.md with the date they were read — read them there, do not put them on the screen.
- **No recommendation and no ranking.** Faculty in the room have institutional constraints you do not know about. The self-test enforces this one mechanically.
- **No claim that going inward is better.** If your delivery accidentally implies it, row 4 is the correction and you should let the page make it.
- **No screenshots.** They would break the offline rule and they rot faster than anything else available.

Before you present

- Open **Settings** → **Presentation mode** for projector-sized type.
- Run the self-test once so you know it is green before anyone asks.
- **Re-read Act III the morning of.** This is the one demo in the collection with a real shelf life. If a name on screen has changed, say so from the stage — it is a better moment than a correct page would have been.
- Reloading resets everything; the app saves nothing anywhere.
- No network is needed at any point. The source links need one, but only if somebody asks you to follow one, and you can promise to send it instead.

Front Doors — Bryant Harrison, Murray State University. Product claims checked 8 September 2026 against first-party sources. Every claim on screen carries a confidence label and links to its origin; everything researched and left off is listed in Act III under “Researched, and deliberately not on the page”.