

PROMPTING STRATEGIES

THE STRANGER · PRESENTER SHEET · 12 MINUTES

LEARNING GOALS

The three stage Takeaways, verbatim from the app. Each is a plain claim the activity demonstrably shows, and each is what that stage's check-yourself item tests.

1. **CONTEXT** With no context the answer is aimed at everybody, which is another way of saying nobody. Context changes **who it is for** — not its length, and not its vocabulary.
2. **SUCCESS SPEC** An unspecified target does not stay unspecified — something else fills it in, and does not tell you. Saying what finished looks like takes that decision back.
3. **EXPERT KEYWORDS** Asking in a field's vocabulary gets that vocabulary back — words nobody in the room supplied. That is real. It is also not the same thing as a better answer, and telling those apart is the skill.

WHERE THIS SITS ON THE PATH

Demo 5 of 17 · Part three. Demo 2 showed that a model reads its context; here the room controls that context, one lever at a time. It hands forward to Demo 6: a perfect prompt in a browser tab still cannot touch your files, so the surface decides reach.

ONE QUESTION FOR THE SESSION

“Who decided what a good answer to your question looks like?”

THE ROLEPLAY — THREE ACTS

You are the Stranger. A volunteer is the Asker. The app is on the projector behind you; run the same states on it as you say them out loud.

ACT I — CONTEXT

~3 min

“I am a stranger on a street corner. I am not stupid. I know a great deal about almost everything. I have never met you.”

1. The Asker says only: “Help me get ready for my test.” Answer out loud, as blandly and generally as you honestly can. Cover every case. Do not clown it — a real assistant is competent here, just aimed at nobody.
2. Ask the room: what would I have needed to know? Take three answers, put them on the board.
3. The Asker reads the context block. Answer again, aimed at that kid.

“I didn't get smarter. You got specific.”

ACT III — THE FOLDERS

~4 min

Three labelled folders or sheets: **BIOLOGY, BRIDGES, WRITING AT WORK**. “I have a folder on nearly everything. The words you use decide which one I open.”

1. The Asker: “Why do bridges stay up?” Open nothing. Answer in plain words: the bottom gets stretched, the top gets squished. All true. All unusable.
2. The Asker adds the keyword block. Open BRIDGES and answer in the real vocabulary — load path, bending moment, neutral axis, section depth.
3. On the app, press H. Count aloud the expert words in the answer that the Asker never typed.

“You didn't know those words. You didn't have to. You only had to point at the folder.”

Say this before you move on — it matters. “There are no folders. I don't open anything. Your words change which words are likely to come next — that is Demo 2's machine, doing this.” The folders are a memory aid for the room. Left unsaid, the prop teaches a filing-cabinet model of a language model, which contradicts Demo 2 — this demo's own prerequisite — and a learner who does Demo 2 then Demo 5 is handed two incompatible machines.

HOUSE RULES

1. **The Stranger is never stupid.** In this demo the bare answers are bland, not broken — we wrote them that way, because the demo dies the moment the bad answer looks rigged. Say it as a rule about the demo, not about models: bare prompts in the wild also produce confident errors, wrong lengths and refusals.
2. **No trick answers.** If the prompt supports a good answer, give a good answer.
3. **Every answer gets a number.** 1–10, from the room, out loud, on the board before the next one.
4. **Nobody says “the AI is bad at this”** without naming which lever is off.
5. **The Asker types nothing they wouldn't really type.** No fair inventing a great prompt to prove a point.

RUNNING ORDER

0:00	Cold open. Act I, live, no screen.
3:00	App, scenario 1, all levers off. Draw the card. Score it.
4:30	One lever at a time. Score each. The curve argues for you.
7:00	Act III on the bridge scenario. Highlighter on, count the words.
9:30	“Which lever would fix this?” — two rounds, hands up.
11:00	Closing slide. Read the line. Stop talking.

PER STAGE — QUESTION, PREDICT PROMPT, EXPECTED OBSERVATION, ANSWER KEY, TEACHING LIMIT

1. **Question:** What changes in the answer when it knows who you are?
Predict: adding context changes *one* dimension — which one? (length / who it is aimed at / which words) The room commits before you press Ask.
Expected observation: the answer becomes aimed at one

ACT II — THE TARGET DECK

~4 min

Hold the eight cards from page two face down. “Somebody has to decide what a good answer looks like. If you don’t decide, I do. And I never tell you what I picked.”

1. Same question, context on, no success spec. **Draw a card and do not show it.** Answer to the card — if it says exhaustive, talk for forty-five seconds without stopping.
2. Now show the card. Wait for the room.
3. Draw a different card. Same question, answer again. Same asker, same words, different answer — and the difference was never theirs.
4. The Asker reads the success spec. Put the deck down and slide it across the table, away from you. On the app, the deck lives *inside* lever 2: flip it on and the three card names get struck out where they sit. Same gesture, same panel.

“You just took the deck out of my hands.”

ON THE APP

? is the **Guide**: what every control does. It opens itself on load. ⚙ beside it is Settings — **Open Presenter Notes** (this sheet, on screen; or press **N**), **Presentation mode** (the projector build: big type, the 1–10 score strip, compare, and the card picker inside lever 2) and **Reset**.

Reset is a fresh load without the reload: question 1, levers off, no answer on screen, scoreboard cleared. It deliberately leaves Presentation mode where you set it, so resetting between classes never drops you into student mode.

asker. Length and vocabulary are not what this lever controls — they follow whichever card is dealt, so they can swing between runs.

Answer key: *who it is aimed at*. The check item's correct option is **Context**; a room that picks Success spec is told the spec is already doing its job.

Teaching limit: a bare prompt is not information-free. “Help me get ready for my test” already leaks topic, register and age. Do not claim it knows nothing.

- 2 **Question:** If you never say what a good answer looks like, who decides instead?

Predict: lever 2 is off, so the target is unspecified — what happens to it? (nothing / something else fills it in / it asks you)

Expected observation: a target card is dealt, face up, and the answer is written to it. Turn lever 2 on and the deck is struck out inside the lever.

Answer key: *something else fills it in*. Correct option **Success spec**.

Teaching limit: the deck is a teaching device, not a mechanism. A real model does not draw from three named cards; its defaults come from training, are more stable than a card draw, and vary continuously. Say so — it is in the Details drawer too.

- 3 **Question:** What do you actually get for asking in the vocabulary of the field?

Predict: keywords off gives at most four expert terms — how many with keywords on? (under five / about ten / fifteen or more)

Expected observation: 15–20 tagged terms, most of them absent from the prompt. Press **H** and count them aloud.

Answer key: *fifteen or more*. Correct option **Expert keywords**. The stage also carries the accuracy-brief case (below), whose correct option is *the prompt only asked for one kind of finding*.

Teaching limit: the counter measures **vocabulary**, not **understanding**, and the big number is a readout of the author's tag list, not a detector. Ask the room the demo's own question: did the answer get more useful, or just more impressive?

MISCONCEPTIONS THIS DEMO CAN LEAVE — NAME THEM, REFUTE THEM

“I watched the model get better as the prompt got better.”

No model ran. All 64 answers were written by hand before the file was saved, and the panel carries a permanent

Scripted

label saying so. What the room watched was an author's four answers, chosen to make one lever's effect legible. The claim the demo can honestly make is *what a better prompt tends to buy you*, never *what this model did*.

“The rising curve proves more levers make better answers.”

It cannot. The answers were written to improve as levers are pulled, the room scores them, and the mean rises — so the curve measures the room agreeing with the author and could not come out any other way. Keep it: making a room commit out loud is good teaching. Do not offer it as evidence, and say which it is.

“If it is not in the prompt box, the AI does not know it.”

False three ways. What it knows *in general* came from training and is not in the box; almost every real product prepends a system prompt the user never sees; and history, retrieved documents and tool output enter the context invisibly. The true and useful version: nothing about *you and your situation* reaches it unless you put it there.

“A bad answer means I left a lever unpulled.”

Sometimes. A lever fixes vague, mis-aimed or wrong-shaped. It does not fix confidently wrong, out of date, refusing, or past what the system can do — and on those, “write a better prompt” is the wrong response. The closing line says both halves; do not shorten it to the first.

“A careful prompt cannot mislead.”

The accuracy-brief case on stage 3 is a documented counter-example from building this very collection: a brief asked for 100% accuracy, every claim that came back was true, and the overall picture was still wrong — because the instruction only asked for one direction of finding. Symmetry is the repair, not more care.

DISCUSSION QUESTIONS

1. The first answer was accurate. Why was it useless?
2. Who decided what a good answer looked like, and when?
3. The expert answer used words you never typed. Where did they come from, and how did asking in plain words keep them hidden?
4. Which of the three levers is hardest to pull for a subject you know nothing about? What do you do then?
5. Name a time you blamed a tool for an answer you had underspecified.
6. Success spec means saying what “done” looks like. When is that hard to say — and what does that tell you about the task itself?

The app is one HTML file and touches no network. Every answer in it was written by hand before the file was saved: there is no model behind the glass. Keys: **1 2 3** levers · **Enter** ask · **Space** skip typing · **H** highlight · **R** replay card · **P** pin to compare. The stage tabs switch levers; **Details** on each stage carries that stage's assumptions, sources and boundary. **Settings** holds **Open Presenter Notes**, **Presentation mode** (projector type plus the class score strip, compare, scoreboard and the card picker) and **Reset**, which restores the demo in place without reloading the page.

The target deck — cut along the dashed lines

Eight editorial choices somebody has to make. When the Asker doesn't make them, the machine does, silently. Hold these face down; draw one whenever the success-spec lever is off.



Length: exhaustive

Cover everything that could possibly be relevant. Leave nothing out.

Turns up in: test · email



Audience: someone who missed the whole semester

Assume they know nothing. Define every term before using it.

Turns up in: test



Format: dense paragraphs

No headings. No bullets. Flowing prose, start to finish.

Turns up in: test · essay



Length: short and safe

Keep it brief. Don't risk being wrong about the details.

Turns up in: bridge



Format: numbered listicle

Everything as a numbered list. Ten items. Equal weight to each.

Turns up in: bridge · email



Audience: an expert peer

Assume deep familiarity. Skip the basics, get to the interesting part.

Turns up in: bridge · essay



Tone: encouraging coach

Lead with affirmation. Soften every critique. Protect motivation.

Turns up in: essay



Coverage: hedge every claim

Present all sides. Avoid committing to one recommendation.

Turns up in: email