

AI Winters: Boom and Bust — presenter guide

Part two of “AI: From Zero to Takeoff”. Runs 12–18 minutes full, 6 minutes cut. One HTML file, offline, no AI at runtime.

Why this demo exists

It follows **Takeoff**, which ends with the room considering AI gains alongside limits and measurement conditions. The honest next question — and you will get it, usually from someone who has been in a technical field a long time — is “*is this another bubble?*”

This demo answers that with evidence instead of reassurance. AI has contracted twice, through different combinations of technical limits, institutional choices, and markets. The histories rhyme; they do not provide a fixed sequence for the present.

*The one sentence to land: **a prediction needs a deadline and an outcome; an assessment needs evidence about its own moment.** Do not grade a historical claim until you know which kind it is.*

The shape of the session

ACT	WHAT HAPPENS	TIME
I — Guess the year	Five time-bounded forecasts, four assessments or untimed ambitions, and one warning, with the speaker and date hidden. Score forecasts; classify the unscored context.	7–9 min
II — The two winters	Scrub the timeline. Five lenses compare historical pathways without forcing a sequence.	3–4 min
III — Anatomy	Compare the two winters and the present under the same lenses. Two current questions stay open.	2–3 min
IV — Now	Rhymes against differences. Close on the checklist.	2–3 min

Act I — running the prediction round

Do the first two cards yourself, out loud, so the room sees the mechanic. Then hand it over — ask for a show of hands on the year before you drag the slider, and ask for a verbal verdict before you press **Lock it in**. The room being wrong together is the experience; a silent room reading a screen is not.

The deck, and what each card is for

YEAR	CARD	TYPE	WHAT IT IS FOR
1950	Turing — fifty years to the imitation game	OPEN FORECAST	Its about-2000 bar was missed; the current status remains contested. Use it to separate a past deadline from an unresolved present judgment.

1958	"walk, talk, see, write, reproduce itself"	UNTIMED AMBITION	An expansive press claim without a deadline. Classify it; do not pretend it supplies a scoreable outcome window.
1957	Simon & Newell — chess champion in ten years	PROMISE	The most useful card in the deck. Right destination, wrong by a factor of four. Pause here.
1960	Simon — any work a man can do, in twenty years	PROMISE	Twenty years lands in 1980, the bottom of the first winter. Also: it is usually misdated to 1965.
1966	ALPAC — machine translation is not near	ASSESSMENT	Distinguish a near-term economic assessment from a forecast. The report also recommended continued computational-linguistics research.
1967	Minsky — substantially solved within a generation	PROMISE	A reported quotation attributed to his textbook. Use it to explain why the app labels this source route reported.
1969	<i>Perceptrons</i> — neural nets read as a dead end	ASSESSMENT	The proofs were right; the history of how the field received them is disputed.
1973	Lighthill — no major impact so far	ASSESSMENT	A retrospective judgment on impact and constraints, not a permanent forecast.
1984	Schank & Minsky — a winter is coming	WARNING	The card that saves the demo from being anti-hype. They were right , three years early.
2024	Amodei — "could come as early as 2026"	OPEN	Deliberately unscored. The year named is the year you are standing in.

Short of time?

Keep four: **Simon & Newell 1957** (over-promise, right destination), **ALPAC 1966** (assessment versus forecast), **Schank & Minsky 1984** (a prediction that was right), and **Amodei 2024** (still open). Those four carry the whole argument. Skip to Act III and close.

What the scorecard is for

It scores time-bounded forecasts and displays dated assessments without grading their future accuracy. That distinction is the teaching point: a claim about conditions in 1966 needs different evidence from a claim about what will happen by 1976.

"Before calling a historical claim right or wrong, ask whether it actually made a prediction. If it did, name the outcome and deadline. If it did not, assess its own evidence."

Act II — the timeline

Do not narrate a cycle into existence. Drag the slider slowly from 1950 and use the rows as comparison lenses: claims, institutional bets, constraints, reassessment, withdrawal. Ask which causes matter in each episode and which do not carry across.

- The ‘ and ’ buttons step one event at a time if you want to stop and read one.
- Tap any dot for the full text and its source.
- Point at the current entries as evidence about today, not a third historical run.

Do not let the lenses turn into a prophecy. The correct line is: “these comparisons tell us what to investigate. They do not establish that a withdrawal will occur.”

The money panel — read the caveat aloud

Four funding examples with different bases: programme budgets, programme expenditure, and a private funding round. They use three currencies and different price years. No continuous, sourceable funding series exists across these years, so the demo shows dated evidence rather than drawing a combined total or a line it cannot defend.

Act III — anatomy

Compare the columns by cause, not by step number. Two cells in the “Now” column remain open: no single binding technical constraint is established here, and this snapshot does not establish a comparable broad contraction.

“Hindsight makes old constraints easier to see. It does not prove that today has the same constraint or that it must end the same way. Ask what evidence would support the claim.”

Act IV — closing

Give both columns roughly equal time. The strongest item on each side:

- **Rhymes:** in February 2026 OpenAI stopped reporting SWE-bench Verified after contamination concerns and an audit of 138 selected hard tasks found faulty tests in 59.4% of that subset. That is a warning about one evaluation, not a field verdict.
- **Different:** scaling behaviour measured across more than seven orders of magnitude, and reported consumer reach at a far larger scale. Earlier booms also had deployed systems and paying organisations, so be precise about the difference.

Close on the checklist, not on a verdict: given any AI claim, ask whether it is a forecast or an assessment, what evidence supports it, and what evidence would move it.

Questions you will get

“So is it a bubble or not?”

Refuse the binary, in one sentence, then give the useful answer: current funding, reported usage, technical reliability, and returns are distinct questions. Earlier eras also had deployed systems and organisational users, though current consumer reach is much larger. A financial correction and a technological dead end are different claims and need different evidence.

“Aren’t the winter dates made up?”

Partly, and the demo says so. The bands are conventions for broader contractions, not dates stamped by a single event. The 1987 hardware collapse is a useful second-winter anchor; the first has more distributed causes and withdrawals, including but not limited to ALPAC. The ends are also conventional and historians disagree; that is why the bands have soft edges and explanatory notes.

“Where did you get this?”

Open **Settings** → **Run the self-test** in front of them. It re-derives, live, that every claim resolves to a declared source, that no source is cited and never used, and that the counts quoted on screen match the data. Then scroll to Act IV, which lists every source and, separately, everything researched and **cut** for lack of one.

“Why is that quote from a quotation-checking website?”

Because the Simon and Minsky lines are routed through a quotation-checking source that cites printed books and page references. This build has not inspected those original pages, so it labels the claims **reported** rather than primary. That distinction is part of the lesson: a plausible citation route is not the same as direct verification.

Before you present

- Turn on **Settings** → **Presentation mode** for projector-sized type.
- Run the self-test once so you know it is green before anyone asks.
- Between rooms, **Settings** → **Reset** returns the whole demo to a fresh load — answers cleared, Act I back at its opening screen, timeline rewind — without a reload. It leaves Presentation mode on, so the projector does not change under you.
- The app saves nothing anywhere, so a reload resets it too.
- No network is needed at any point. If the room’s wifi is down, nothing changes.

The two buttons, top right

- ? — **Guide**: the short how-to panel that opens on first load, reopenable at any point. (Not this document.)
 - ⚙ — **Settings**: **Open Presenter Notes** (this document, on screen), **Presentation mode**, **Reset**, and **Run the self-test**.
-

AI Winters: Boom and Bust — Bryant Harrison, Murray State University. Sources checked September 2026. Every figure on screen links to its origin; everything that could not be sourced is listed in Act IV under “What was cut, and why”.