

Two Winters — presenter guide

Part two of “AI: From Zero to Takeoff”. Runs 12–18 minutes full, 6 minutes cut. One HTML file, offline, no AI at runtime.

Why this demo exists

It follows **Takeoff**, which ends with the room having watched capability curves go almost vertical. The honest next question — and you will get it, usually from someone who has been in a technical field a long time — is “*is this another bubble?*”

This demo answers that with evidence instead of reassurance. AI has collapsed twice. Both times the same five things happened in the same order. Knowing that sequence is how you judge a claim without either swallowing it or dismissing it.

*The one sentence to land: **confident predictions that AI was nearly here have a bad record — and so do confident predictions that it would never work.** Most rooms only believe the first half until they are scored on the second.*

The shape of the session

ACT	WHAT HAPPENS	TIME
I — Guess the year	Ten sourced predictions with the speaker and date hidden. The room guesses the year and the outcome. Scorecard splits promises from dismissals.	7–9 min
II — The two winters	Scrub the timeline. The five lanes light up in order, twice.	3–4 min
III — Anatomy	The two winters and the present, on the same five stages. Two cells in the third column are empty.	2–3 min
IV — Now	Rhymes against differences. Close on the checklist.	2–3 min

Act I — running the prediction round

Do the first two cards yourself, out loud, so the room sees the mechanic. Then hand it over — ask for a show of hands on the year before you drag the slider, and ask for a verbal verdict before you press **Lock it in**. The room being wrong together is the experience; a silent room reading a screen is not.

The deck, and what each card is for

YEAR	CARD	TYPE	WHAT IT IS FOR
------	------	------	----------------

1950	Turing — fifty years to the imitation game	PROMISE	Opens gently. Nobody guesses this early. Sets up that the field has always had deadlines.
1958	"walk, talk, see, write, reproduce itself"	PROMISE	The press-release failure mode. The underlying idea was real and important; the sentence was not.
1957	Simon & Newell — chess champion in ten years	PROMISE	The most useful card in the deck. Right destination, wrong by a factor of four. Pause here.
1960	Simon — any work a man can do, in twenty years	PROMISE	Twenty years lands in 1980, the bottom of the first winter. Also: it is usually misdated to 1965.
1966	ALPAC — machine translation is not near	DISMISSAL	The first dismissal. Fair about 1966, catastrophic as a forecast. Funding stopped.
1967	Minsky — substantially solved within a generation	PROMISE	From his own textbook, not a magazine. Use this if someone challenges the sourcing.
1969	<i>Perceptrons</i> — neural nets read as a dead end	DISMISSAL	The proofs were right; the conclusion the field drew was not. And the authors disputed the tidy story.
1973	Lighthill — no major impact anywhere in the field	DISMISSAL	The report that ended UK funding. True about 1973; read as permanent.
1984	Schank & Minsky — a winter is coming	WARNING	The card that saves the demo from being anti-hype. They were right , three years early.
2024	Amodei — "could come as early as 2026"	OPEN	Deliberately unscored. The year named is the year you are standing in.

Short of time?

Keep four: **Simon & Newell 1957** (over-promise, right destination), **ALPAC 1966** (dismissal that aged badly), **Schank & Minsky 1984** (a prediction that was right), and **Amodei 2024** (still open). Those four carry the whole argument. Skip to Act III and close.

What the scorecard is for

It shows two bars: how the room did on promises, and how it did on dismissals. Most rooms are much better at spotting hype than at spotting a dismissal about to age badly. If the two bars come out very different heights, say so out loud — that asymmetry is the finding, not the total.

"You are all good at not being fooled by someone selling you the future. You are much worse at not being fooled by someone telling you it will not work. Both of those are in the historical record about equally often."

Act II — the timeline

Do not narrate the events. Drag the slider slowly from 1950 and let the room watch the rows fire top to bottom. Say the stage names as each lane lights: promise, money, limit, naming, withdrawal. Then keep

dragging and let them notice it happen again.

- The ‘ and ’ buttons step one event at a time if you want to stop and read one.
- Tap any dot for the full text and its source.
- Point at the **third run**, from 2020 on: the top rows have fired and the bottom two have not.

Do not let the lanes turn into a prophecy. The correct line is: “the first three stages have happened. Whether the last two do is the open question, and nothing on this screen settles it.”

The money panel — read the caveat aloud

Four funding commitments, in four different currencies, not inflation-adjusted. That is on purpose: no continuous, sourceable series of AI funding exists across these years, so the demo shows four points with sources rather than drawing a line it cannot defend. Saying that out loud buys you more credibility than a prettier chart would.

Act III — anatomy

Read down the first column, then the second, then stop at the third. Two cells in the “Now” column are empty — the binding technical limit, and the withdrawal. Make the room sit with that.

“In both previous cycles, the limit was completely obvious afterwards and invisible at the time to the people doing the work. There is no reason to think we are better placed. If someone tells you they know what today’s binding constraint is, they are making the same kind of claim that failed twice.”

Act IV — closing

Give both columns roughly equal time. The strongest item on each side:

- **Rhymes:** in February 2026 OpenAI stopped reporting SWE-bench Verified because models had been trained on it and most of the hard problems they still failed had broken tests. Scores kept rising; what they measured stopped being capability. Both previous winters started with a gap between claim and measurement.
- **Different:** scaling behaviour that has held across more than seven orders of magnitude, and 900 million people a week actually using the thing. Neither previous boom had a measured extrapolable relationship, and neither had users.

Close on the checklist, not on a verdict: given any AI claim, ask which of the five stages it belongs to and what evidence would move it.

Questions you will get

“So is it a bubble or not?”

Refuse the binary, in one sentence, then give the useful answer: some of the money is ahead of the returns, which is stage two of the pattern; and the technology is in hundreds of millions of hands, which neither previous boom managed. Both are true. Financial correction and technological dead end are different claims and they are being argued as if they were one.

“Aren’t the winter dates made up?”

Partly, and the demo says so. The starts are well attested — ALPAC 1966, the hardware collapse in 1987. The ends are conventional and historians disagree; the second winter is dated to 1993 here and to 2000 by other accounts. That is why the bands have a soft right edge and a note under them.

“Where did you get this?”

Open **Settings** → **Run the self-test** in front of them. It re-derives, live, that every claim resolves to a declared source, that no source is cited and never used, and that the counts quoted on screen match the data. Then scroll to Act IV, which lists every source and, separately, everything researched and **cut** for lack of one.

“Why is that quote from a quotation-checking website?”

Because the Simon and Minsky lines are in printed books from 1960 and 1967 that are not online, and that site verified both against page scans and gives the page numbers. It is a stronger source than a paraphrase from a blog, and the demo marks the difference: every claim carries **primary** or **reported**.

Before you present

- Open **Settings** → **Presenter mode** for projector-sized type.
- Run the self-test once so you know it is green before anyone asks.
- Reloading resets everything; the app saves nothing anywhere.
- No network is needed at any point. If the room’s wifi is down, nothing changes.

Two Winters — Bryant Harrison, Murray State University. Sources checked September 2026. Every figure on screen links to its origin; everything that could not be sourced is listed in Act IV under “What was cut, and why”.